

Simulating High Availability Scenarios in Cloud Data Centers : A Closer Look

Mohamed Abu Sharkh*, Abdallah Shami*, Peter Öhlén†, Abdelkader Ouda*, and Ali Kanso†

*Department of Electrical and Computer Engineering, Western University

*London, Canada, Email: {mabusha, ashami2, aouda}@uwo.ca

† Ericsson Research, Stockholm, Sweden, Email: {ali.kanso, peter.ohlen}@ericsson.com

Abstract—Migrating to the cloud is becoming a necessity for the majority of businesses. Cloud tenants require certain levels of performance in aspects like high availability and service rate and deployment options. On the other hand, Cloud providers are in constant pursuit of a system that satisfies client demands for resources, maximizes availability, minimizes power consumption and, in turn, minimizes the cloud providers' cost. A main challenge cloud providers face here is ensuring high availability (HA). High availability includes the combined reliability of components of all categories including network, computational, hardware and software components of all layers.

In this work, we first address the need for a cloud simulator that enables HA algorithm testing in cloud environments and observe its impact on energy efficiency. We introduce a framework to amend cloud simulators with critical HA features. We take GreenCloud, a major simulator with a direct focus on green computing, and implement these features as an additional measurement layer. We demonstrate these added features by simulating their impact on a phased communication application (PCA).

Keywords: *Cloud Computing, Cloud simulators, High Availability, Scalability, Virtualization, Network and systems monitoring and measurements*

I. INTRODUCTION

With the increased migration of business applications to the cloud, more familiar challenges related to service performance are arising. The Cloud tenants require certain levels of performance in aspects like high availability, pricing options and general reliability. On the other hand, Cloud providers strive to satisfy their clients' demands according, first, to their resource requirements and, second, to their service quality requirements. Cloud providers are in constant pursuit of the holy grail of cloud management systems. This translates to a system that satisfies client demands for resources, maximizes availability, minimizes power consumption and, in turn, minimizes cloud providers' cost.

A main challenge cloud providers face to achieve these goals is ensuring high availability (HA). Client are becoming more demanding in that aspect and the 5 nines requirement (guaranteeing that the data center service is available for 99.999% of the time) is becoming a reality. High availability includes the combined reliability for components of all categories including hardware and software components. This

covers Network and processing resource sides. It also includes the availability of components of all layers (cloud, racks, servers, VMs, applications, application components).

HA received major attention as soon as the cloud solutions started being deployed. In [1] for example, we notice a solution depending on RAID-technology used to accomplish the task of managing data across multiple cloud storage providers. Factors like geographic location, quality of service, providers' reputation, and budget preferences are taken into consideration. Early efforts up to 2012 are summarized in [2] which focused on HA techniques based on multi-core processing, virtualization, and distributed storage. In more recent work, the authors of [3] presented the HA constraints within the context of a multiple objective resource scheduling problem in the cloud environment. Their algorithm was tested on a real life social news application with synthetically generated costs and loads. Another variation is the one seen in [4], where a failover strategy is presented. The tested technique combines load balancing algorithms with multilevel checkpointing so as to decrease checkpointing overheads. In their paper published in 2014, the authors of [5] offer an architecture for automatic failover between multiple Platform-as-a-Service (PaaS) cloud providers.

Virtualized storage and redundancy are of prime interest as well. In addition, other works can be found discussing experiments on VM migration analysis, including security related aspects [6] and performance analysis of migration algorithms [7]. Finally, in [8], a solution is provided to ensure HA in cloud storage by decreasing virtual machine reboot time.

Cloud simulators play a critical role while developing the optimal cloud data center management System. A cloud simulator offers an environment to implement scheduling policies and operation scenarios with clear cost advantage and decreased risk. A cloud simulator serves as the first barrier that can examine resource allocation algorithms, energy efficiency techniques and HA-aware scheduling algorithms. A comprehensive survey and feature comparison of the available cloud simulators can be found in [10].

The most popular simulators tried to cover cloud functionality in a generic view that covers components, processing components and data center management as can be seen in efforts like CloudSim (along with the extensions) [11][12],

¹This work is supported by the Natural Sciences and Engineering Research Council of Canada (NSERC-STPGP 447230) and Ericsson Research.




Fig. 1. GreenCloud output

GreenCloud [13] and MDCsim [14]. With all the efforts to propose HA solutions for the different elements composing the cloud, we still lack a Cloud simulator that can simulate the behavior of multi-tiered applications, while considering the different failures that can occur in the Cloud, and quantify their impact on the applications availability.

Therefore, there is a need for a cloud simulator that enables high availability algorithm testing in order to reach a HA scheduling technique that does not sacrifice energy efficiency. To the best of our knowledge, there is not a simulator that provides the detailed functionality that enables measuring HA metrics, testing HA algorithms and producing the results in a way that can serve academia and the industry.

In this work, we introduce a framework to amend cloud simulators with HA features. We take GreenCloud as an example of a major simulator with a direct focus on green computing and implement these features as an additional measurement layer. This is illustrated using the specifications of a phased communication application (abbreviated henceforth as PCA).

Section II introduces a framework to enhance the cloud simulator of choice with required features in order to turn it into an HA-aware simulator. Section III presents the enhanced GreenCloud architecture and added features. Section IV explains the experimental setup. Section V shows some of the testing results achieved after enhancing GreenCloud with HA-awareness and then we present conclusions in section VI.

II. A FRAMEWORK TO IMPLEMENT HA-AWARENESS IN CLOUD SIMULATORS

The notion of guaranteeing a certain standard of high availability has a direct effect on performance and specially energy consumption in a cloud environment. Implementing redundancy which is a prerequisite to any highly available algorithm causes the amount of resources required to serve a certain load to increase linearly. The increase factor in this case will depend on the number of redundant hardware components. In addition, the activities required to implement redundancy

require computational and network resources just like user requests. This will, in turn, increase the demand for data center resources and naturally increase energy consumption. Hence, a careful consideration of the effect HA algorithms have on the energy efficiency in the cloud is needed. This will assist cloud providers in making the decision as to which HA algorithm achieves the best trade-off for them in terms of availability and energy efficiency.

To enable any cloud simulator to measure availability and evaluate different HA algorithms in a cloud data center, it has to include the following features.

1) HA Features

a) Component Failure

Components failure simulation means that simulator users would be able to inject failure events based on any specific time series or distribution that suits their input. Components covered should ideally include all GreenCloud components. This includes servers, racks, virtual machines, switches (all types) and other network components. This includes the case of the whole data center failing in case of a major power problem for example. Another issue here is the domino effect of failure. If a host fails, all VMs on it and tasks scheduled on them would have to suffer failure too. All these resource consumers (as termed in GreenCloud) would be rescheduled on other servers(resource providers).

b) Component Recovery

In a real data center, failed components naturally come back to a running state after a period of time.

c) Synchronization Redundancy and Server Groups

Based on the roles they can perform, servers are divided into groups. A task in phase I is only scheduled on a server in the corresponding(matching) group. As shown in Fig. 2 and the flowchart in Fig. 3, as the task in the first phase (for instance, PCA client phase) is executed, the following task is added to the scheduling queue matching the next phase(PCA). When that phase is over, the following task is added to the following queue(IMS-CSCF) until the last phase. That simulator also guarantees that if a task in the IMS-CSCF phase fails for any reason(server failure or VM failure for example), the task will only be scheduled on a server from the same server group(IMS-CSCF server group).

2) Workload Modeling and Scheduling

a) Higher Granularity Modeling of Application

i) Adding User-defined Resources

A task as defined in GreenCloud includes the attributes seen in Table I. The ability to add extra types of resources means more shades of problems can be represented. We can represent scenarios in which specific servers have re-

sources other servers do not. For example, some servers would be able to offer DB service for a specific number of requests per second. This can be helpful in simulating scenarios where each server group performs a predefined role.

ii) Defining Tasks with User-defined Resource Requirements

Emphasizing on the previous feature, task resource specification has to be amended with the amount the task would request of the new resource (I.O percentage or number of HTTP requests the task contains for example).

iii) Defining Applications with Diverse Set of Tasks Representing Execution in Dynamically-defined Phases

A request in this case moves from being just a simple task disjoint from all other tasks and defined by only its computational resource requirements and data exchange requirements to something more inclusive. A request consists of a set of tagged tasks that have functional dependency between them. The first task would represent the first phase in implementing the complete request. The request would be fully executed when the task representing the last phase gets executed. As seen in the figure , the request consists of tasks being processed at the App front end component, app core , app back end then the app core again and the front end component once more before eventually sending the results to the client. Each one of these components could be scheduled on a distant VM on a separate server along the data center or event in other data centers.

b) Functional Dependency and Dynamic Scheduling of Phases (based on resource requests)

Once the request is constructed of a set of tasks that differ in resource requirements, the scheduler is supposed to take that into consideration while scheduling. A task will only be scheduled on a server with the sufficient resources and also the server with the suitable functionality. Tasks in phase i will not be processed or scheduled before tasks in phase i-1 are completed.

3) Monitoring and Reporting

a) Reporting Availability Status The measurements shown in figure 1 are something to start with. The metrics included are:

- Failure cases/component (server, VM, etc)
- Failure events/task
- Total outage time and percentage.

The availability of a complex component like a data center can be calculated factoring in the availability of its subcomponents (server, racks, switches , etc).

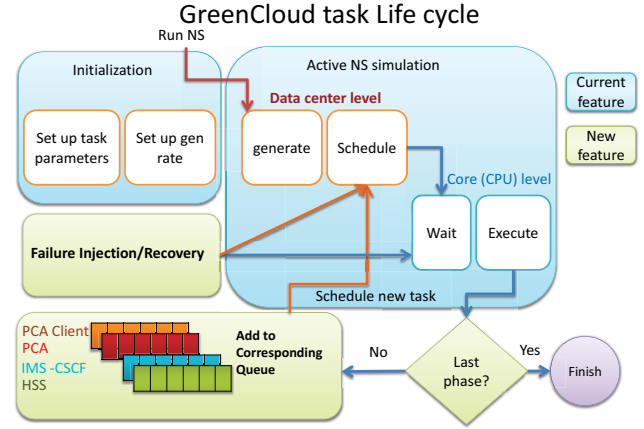


Fig. 2. The enhanced task life cycle in GreenCloud

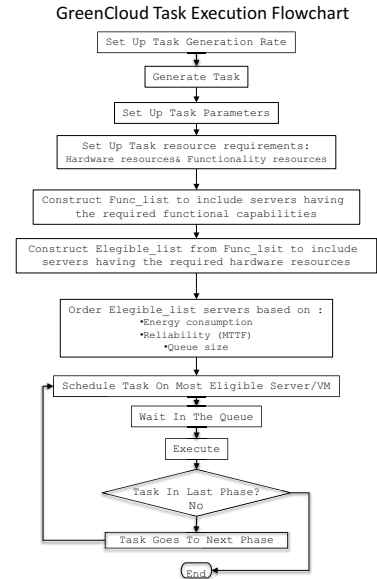


Fig. 3. GreenCloud task execution flowchart

b) Reporting Service Total Exec Time & Outage Time

The average completion time of a full request is shown in Fig. 1. This means the total time to complete all of its tasks (phases).

III. ENHANCING GREENCLOUD WITH HA FEATURES

A. GreenCloud scheduling environment

As one of the most detailed cloud simulators available, Greencloud arises as a powerful tool to implement the proposed functionalities. GreenCloud was developed as simulator with a focus on energy efficiency and fine grained networking capabilities. The prime purpose cited for building GreenCloud is mitigating overprovision issues [13]. Overprovision happens in a data center due to the loads constantly changing on

TABLE I
TASKS IN GREENCLOUD- DEFINING ATTRIBUTES

Variable name	Attribute
task(size)	Input data to be sent to host the task is scheduled on
task(memory)	[Byte] of used RAM
task(storage)	[Byte] of disk space
task(duration)	computing deadline of tasks in seconds(can be set as parameter)
task(duration)	computing deadline in seconds
task(outputsize)	standard- Size of output on task completion
task(outputsize)	Size of output on task completion
task(outputsize)	low comm- Size of output on task completion
task(intercom)	Size of inter-task communication

the computational and communication resources. The average load can be as low as 30% of the data center server and network capacity [13]. This, in turn, causes the data center to systematically use more power than the optimal value.

GreenCloud offers simulation capabilities including multiple topology choices (2 layers and 3 layers) and it offers communication through packets using the underlying NS-2 simulator features. GreenCloud also offers the choice of scheduling tasks(user requests) on hosts directly or on virtual machines which reside on hosts.

Tasks are modeled as unit requests that contain resource specification in the form of computational resource requirements (MIPs, memory and storage) in addition to data exchange requirements (task size variable representing the process files to be sent to the host the task scheduled on before execution, data sent to other servers during execution and output data sent after execution).

B. HA enhancements made to GreenCloud

We have implemented the HA features discussed in the previous section in GreenCloud. As illustrated in Fig. 1, failure injection feature was added to GreenCloud. Failures and recovery can be injected at preset times or taken from a file. Related recovery procedures were also added. Corresponding service and outage metrics are also available in the enhanced GreenCloud as seen in the figure. Functional dependency and synchronization dependency were implemented including implementing task tagging, server grouping and the related service metrics. Major changes to the GreenCloud Scheduler are illustrated in the flowchart in Fig. 3. For a start, we distinguish between hardware resource like RAM, CPU, disk, bandwidth and Functionality resource like the HTTP capabilities, the availability of DB implementation on the server and so on.,

The scheduler starts by filtering out the list of the all servers, leaving only those that do include the functionality resource producing a Func_list. Next, the servers that do not satisfy the hardware resources from the Func_list are filtered out and hence an Eligible_list is produced. This list is ordered based on a combination of consumption of energy, reliability (MTTF) and queue size. The task is then scheduled on the most eligible server. This is done in every phase the task is in until the whole request is served.

In the following paragraphs, the scenario used to test these features is illustrated.

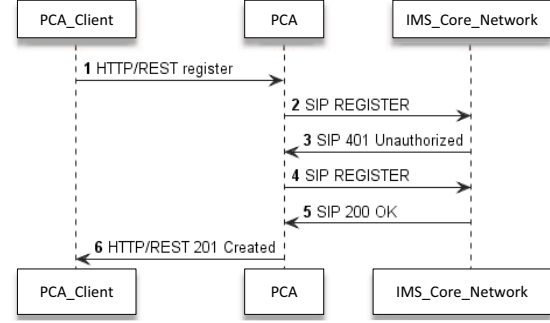


Fig. 4. Sequence diagram of the PCA scenario-Register

C. Phased Communication application(PCA) Scenario

The scenario we chose to evaluate the framework through is a case where the cloud client uses leased cloud resources to simulate the implementation of Phased Communication application(PCA). This application makes communication services accessible to developers using technologies like HTTP and WebRTC. Without loss of generality, this scenario is used as it demonstrates functional dependency, synchronization dependency and the tagged tasks phases.

1) *Registration*: In this scenario, a Web Access component allows HTTP-based clients to connect to an IMS network. The PCA receives HTTP/REST-based requests from a web client and sends these requests to the IMS network. The PCA receives requests from the IMS network and sends them as events over an event channel to the web client. these requests allow clients to register to the IMS Core Network. The resource requirements for requests in this phase are as follows. (i) CPU: 1200 registrations per second where each registrations consumes 2 MIPS; (ii) Memory: 1 active registration, uses 13905 bytes of RAM for the duration of the registration; (iii) Network: 1 HTTP POST per registration (TCP); 2 SIP Register per registration (UDP); (iv) Duration of the registration is assumed to be 24 hours.

Once a registration is successful, PCA will automatically re-register the request at regular interval 3600 seconds (default).Re-registration consumes very little RAM.The resource requirements for requests in this phase are as follows. (i)CPU: 2400 re-registrations, each re-registration uses 1.625 MIP (ii)Network: 2 SIP Register per registration (UDP) This scenario is illustrated in the sequence diagram in Fig. 4.

2) *Audio-Video Calls*: A registered PCA Web Client can initiate an audio/video call to a remote user and the call is accepted/started. The resource requirements for requests in this phase are as follows. (i) CPU: 1000 calls per second where each call consumes 2.4 MIPS; (ii) Memory: 1 active call, uses 20210 bytes of RAM for the duration of the call; (iii) Network:2 HTTP POST per call (TCP); 2 HTTP GET;1 SIP INVITE (UDP);1 SIP ACK (UDP); 1 SIP BYE (UDP); (iv) Call Mean Hold Time is assumed to be 180 seconds. This scenario is illustrated in the sequence diagram in Fig. 5.

When implementing this scenario in the enhanced Green-

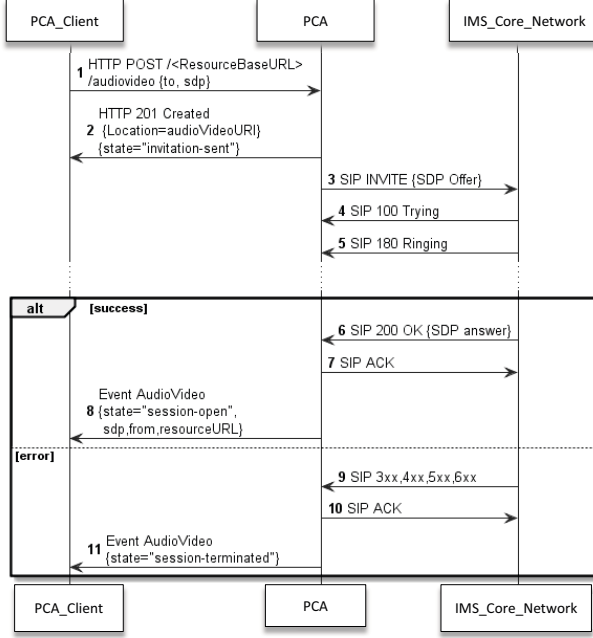


Fig. 5. Sequence diagram of the PCA scenario-Audio-Video Call

TABLE II
SIMULATION PARAMETERS

Parameter	Value/calculation Method
Server count	tested for 144 and 1536 servers
Topology	3-tier topology
CPU configuration	HP ProLiant BL460c Gen8 Server Blade
Problem size	10 up to 100 requests/second
Average server outage time	$\sum_{i=1}^n$ (average outage for servers in $group_i$)
Average service completion time	$\sum_{i=1}^n$ (average execution time for tasks in $phase_i$)
Total simulated time	up to 3000 time units(seconds)
Average Experiment time	1-2 hours

Cloud, phases of the full request can be defined by determining the steps where messages are exchanged and processing is required by the PCA and the components interacting with it.

For example, in the registration scenario, a request would go through the following phases and each phase would yield a task to be scheduled and served.

$PCA - App \rightarrow PCA \rightarrow IMS - CoreNet$.

In the next section, initial testing results of the HA enhanced GreenCloud are presented.

IV. INITIAL RESULTS

Major simulation parameters used include typical data center configuration parameters like the ones in table II. server count, topology and server resource configuration all go under that category. Simulator and input configuration parameters like input size, request resource specifications and total simulated time are also important here. In addition, the methods of calculating average service time and average outage time are also documented in table III.

As mentioned in the previous sections, we used the enhanced GreenCloud to test the implemented HA features and in order to implement PCA scenario which spans all these features. Moreover, we show the effect of some of the parameters on major metrics like the request average response time.

First, Table III shows some of the results for testing the PCA scenario while changing the problem size (represented by the number of request/second). The obvious trend here is an increase in the average response time as the problem size (arrival rate) grows. We used minimal number of servers in order to stress the system as much as possible in this scenario.

Table IV shows the impact of increasing the dequeuing rate of the phased request. As the request finished the first phase it goes into the scheduling queue to be executed for the following phase. this queue is dequeued with the rate shown in the table. As the rate grows the average execution time (response) for a request decreases of up to 37%. It keeps decreasing until a limit where it stabilizes and then starts increasing the average response time. This happens because the dequeuing process adds no more value and just causes an overhead for the simulator. the simulator runs the dequeuing process while the queue is empty.

Table V show the effect of changing the server group sizes on the average response time of a request. To show that and also to show that this scenario works for large data center, we tested the same set of features for a large data center of 1536 (the three-tier topology configuration in GreenCloud). A gain of up to 17% in the average response time can be reached by controlling the server group sizes even for the very small load of 10 requests per second. The lesson to be learned here is that the parameter configuration in terms of server group distribution and dequeuing rate can have positive impact if they were adjusted to the exact demands of the input request set. A closer look at the resource requirements and life time of each phase of the request has the potential to produce sizeable gains in response time and power consumption. An effect on the power consumption metric was recorded as well. For the large data center setup (1500 servers), A gap of 14.5 kW*h was recorded between the two server group configurations used. This gap could grow as the load starts increasing to meet the data center capacity. The lesson to be learned here is that the parameter configuration in terms of server group distribution and dequeuing rate can have positive impact if they were adjusted to the exact demands of the input request set. A closer look at the resource requirements and life time of each phase of the request has the potential to produce sizable gains in response time and power consumption.

V. CONCLUSION

The cloud computing environment current state of affairs imposes the need for cloud solutions that enables high availability capabilities without sacrificing energy efficiency. To address this need, researchers need resilient comprehensive algorithms as well as sufficient tools that enables them to

TABLE III
PCA SCENARIO SIMULATION RESULTS- IMPACT OF INCREASING THE
ARRIVAL RATE ON AVERAGE SERVICE TIME

Dequeuing Rate	Total servers	Group 1 size	Group 2 size	Requests per sec	Average service time
1 call/sec	2	1	1	10	0.36190456
1 call/sec	2	1	1	50	0.74810514
1 call/sec	2	1	1	100	1.24199910

TABLE IV
PCA SCENARIO SIMULATION RESULTS-IMPACT OF THE DEQUEUING RATE
ON THE AVERAGE SERVICE TIME

Dequeuing Rate	Total servers	Group 1 size	Group 2 size	Requests per sec	Average service time
1 call/sec	140	70	70	100	1.22678165
4 calls/sec	140	70	70	100	0.84924532
1000 calls/sec	140	70	70	100	0.77226364
1 call/sec	2	1	1	100	1.24199910
4 calls/sec	2	1	1	100	0.89261223
1000 calls/sec	2	1	1	100	0.79194872

evaluate new techniques. We introduced a framework to amend GreenCloud cloud simulator with HA features. The sides covered were HA features, workload modeling and scheduling features, and reporting/monitoring. This was illustrated using the specifications of the PCA example.

After implementing the PCA scenario for different topologies and data center configurations, the results show that gains in the average response time can be achieved by controlling parameters including the server group distributions and dequeuing rate. This includes adjusting the scheduling decisions and the parameter configuration according to the demands of the input request set. A closer look at the resource requirements and life time of each phase of the request has the potential to produce sizable gains in response time and power consumption.

In the future, we will investigate the performance of different scheduling algorithms when amending the resource list with user-defined unconventional resources like security capabilities or guaranteed DB performance. The ultimate objective is reaching a comprehensive HA algorithm that can serve a wide range of cloud applications.

ACKNOWLEDGMENTS

The authors would like to thank GreenCloud developing team represented by Mateusz Guzek for promptly answering questions. The authors would also like to thank the CloudSim team for an informative forum and group website.

REFERENCES

- [1] M. Schnjakin, R. Alnemr, and C. Meinel, "A security and high-availability layer for cloud storage," Grid Computing Environments Workshop, GCE '08, pp.1-10, 12-16 Nov 2008.

TABLE V
PCA SCENARIO SIMULATION RESULTS-LARGE DATA CENTER

Group 1 size	Group 2 size	Average service time	Power consumption
750	750	0.26823299	17111.9 kW*h
1500	1	0.31653402	17095.4 kW*h

- [2] C. Pham, P. Cao, Z. Kalbarczyk, and R. Iyer, "Toward a high availability cloud: Techniques and challenges," IEEE/IFIP 42nd International Conference on Dependable Systems and Networks Workshops (DSN-W), Jun. 2012, pp. 1-6.
- [3] M. Frincu and C. Craciun, "Multi-objective Meta-heuristics for Scheduling Applications with High Availability Requirements and Cost Constraints in Multi-Cloud Environments," IEEE International Conference on Utility and Cloud Computing (UCC), Washington, DC, USA, 2011.
- [4] D. Singh, J. Singh, and A. Chhabra, "High availability of clouds: Failover strategies for cloud computing using integrated checkpointing algorithms," 2012 International Conference on Communication Systems and Network Technologies (CSNT), May 2012, pp. 698703.
- [5] I.D. Addo, S.I. Ahamed, and W.C. Chu, "A Reference Architecture for High-Availability Automatic Failover between PaaS Cloud Providers," International Conference on Trustworthy Systems and their Applications (TSA), 9-10 Jun 2014, pp.4 - 21.
- [6] Y. Wang, J. Ma, D. Lu, X. Lu, and L. Zhang, "From high-availability to collapse: quantitative analysis of Cloud-Droplet-Freezing attack threats to virtual machine migration in cloud computing," Cluster Computing Journal, Dec 2014, Volume 17, Issue 4, pp 1369-1381
- [7] D.K. Prasad, "High Availability Based Migration Analysis to Cloud Computing for High Growth Businesses," International Journal of Computer Networks (IJCN), Volume 4, Issue 2, 2012 pp 35
- [8] S. Jaffer, M. Chitnis, and A. Usgaonkar, "Providing high availability in cloud storage by decreasing virtual machine reboot time," 10th Workshop on Hot Topics in System Dependability (HotDep 14), Oct 2014.
- [9] M. Abu Sharkh, M. Jammal, A. Ouda, and A. Shami, "Resource Allocation In A Network-Based Cloud Computing Environment: Design Challenges," Communications Magazine, IEEE Volume:51, Issue:11 Nov., 2013, pp.46-52.
- [10] M. Abu Sharkh, A. Kanso, A. Shami, and P. Öhlén "Building a Cloud on Earth: A Study of Cloud Computing Data Center Simulators," Submitted to Elsevier Computer Networks journal, Aug. 2015.
- [11] R. N. Calheiros, et al., "CloudSim: a toolkit for modeling & simulation of cloud computing environments & evaluation of resource provisioning algorithms," Software: Practice & Experience, Vol.41, pp.23-50, 2011.
- [12] S. K. Garg and R. Buyya, "NetworkCloudSim: modeling parallel applications in cloud simulations," 4th IEEE International Conference on Utility and Cloud Computing, pp.105-113, 2011.
- [13] D. Kliazovich, P. Bouvry, and S. U. Khan, "GreenCloud: a packet-level simulator of energy-aware cloud computing data centers," Journal of Supercomputing, special issue on Green Networks, 2011.
- [14] S. H. Lim, B. Sharma, G. Nam, E. K. Kim, and C. R. Das, "MDC-Sim: a multi-tier data center simulation, platform," IEEE International Conference on Cluster Computing and Workshops, pp.1-9, 2009.
- [15] M. Abu Sharkh, A. Ouda, and A. Shami, "A Resource Scheduling Model for Cloud Computing Data Centers," in Proc. IWCMC, 2013, pp.213-218.
- [16] M. Jammal, A. Kanso, and A. Shami, "High Availability-Aware Optimization Digest for Applications Deployment in Cloud," in Proc. ICC, 2015, in press. Available: <http://vixra.org/pdf/1410.0193v1.pdf>
- [17] M. Jammal, A. Kanso, and A. Shami, "CHASE: Component High-Availability Scheduler in Cloud Computing Environment," in Proc. IEEE International Conference on Cloud Computing (CLOUD), 2015 pp. 477-484.