

An Extensible Architecture for Detecting Violations of a Cloud Environment's Constraints During Legacy Software System Migration

Sören Frey and Wilhelm Hasselbring
 Software Engineering Group
 University of Kiel
 24118 Kiel, Germany
 {sfr, wha}@informatik.uni-kiel.de

Abstract—By utilizing cloud infrastructures or platforms as services, SaaS providers can counter fluctuating loads through smoothly scaling up and down and therefore improve resource- and cost-efficiency, or transfer responsibility for the maintenance of complete underlying software stacks to a cloud provider, for instance. Our model-based approach CloudMIG aims at supporting SaaS providers to semi-automatically migrate legacy software systems to the cloud. Thereby, the analysis of conformance with the specific constraints imposed by a cloud environment candidate along with the detection of constraint violations constitutes an important early phase activity. We present an extensible architecture for describing cloud environments, their corresponding constraints, and appropriate violation detection mechanisms. There exist predefined constraint types with specified domain semantics as well as generic variants for modeling arbitrary constraints. A software system's compliance can be examined with the assistance of so called constraint validators. They operate on discovered KDM-based models of a legacy system. Additional constraint validators can be plugged into the validation process as needed. In this context, we implemented a prototype and modeled the PaaS environment Google App Engine for Java. We report on a quantitative evaluation regarding the detected constraint violations of five open source systems.

Keywords—Cloud computing; Cloud environment constraints; Constraint validation; CloudMIG; Migration to the cloud; KDM

I. INTRODUCTION

Adopting cloud computing technologies is a worthwhile option for many companies considering to modernize existing applications, whole application landscapes, and the accompanying procurement, operation, and maintenance processes. Among others, improved scalability, reliability, and the often-cited cut of capex costs constitute veritable stimuli. Furthermore, cloud computing technologies can facilitate the concentration on core business activities for software service providers. This is valid for in-house providers as well. Software service providers are called SaaS providers in the cloud computing context [1]. We investigate obstacles the SaaS providers face when migrating existing applications to the cloud and thereby build upon infrastructures and platforms delivered as services from IaaS and PaaS cloud providers. While the potential advantages accompanied with a migration might be compelling, there exist numerous difficulties SaaS providers have to overcome before being able to utilize cloud technologies and particularly to leverage a cloud's capabilities. A cloud environment may enable superior scalability, but

many legacy applications' architectures are not designed to exploit the provided elasticity, for instance. Therefore, it is often necessary to restructure an architecture to take full advantage of cloud technologies. But even running an unimproved existing application in the cloud while striving after a minimal change set can be a challenging task. Our model-based approach *CloudMIG* [2] aims at supporting SaaS providers to evaluate different cloud environment candidates from a technical perspective and to assist reengineers in accomplishing restructuring activities during the migration process. For being able to compare competing cloud environment offers, one essential aspect is to take adequacy for the specific legacy architecture and implementation into account. Not every cloud environment is similarly suited for a software system as the cloud environments impose varying constraints on the applications they host. For example, the ability to access the underlying file system may be permitted in a non-uniform way, or the usage of particular network protocols might be restricted. Hence, for each cloud environment being under consideration a detailed analysis has to be performed to validate the specific constraints.

We present an extensible architecture for modeling those cloud environment constraints (*CECs*) and detecting cloud environment constraint violations (*CEC violations*) to simplify the conformance checking process. Thus, the specific *CECs* for a cloud provider can be documented in a reusable manner in a so called cloud profile and serve as a validation input for arbitrary legacy systems. The conformance of a software system with the modeled *CECs* can be examined with the help of constraint validators. Each constraint validator can check an existing or reconstructed model of the software system for code artifacts that would lead to *CEC violations* when being deployed unmodified. Here, every constraint validator refers to a certain constraint type. For cases where the standard constraint validators might not yield sufficient results or more efficient validation processing can be applied, the default set of validators can be extended. To ease the provision of detection mechanisms in such cases, additional constraint validators can be plugged into the overall validation process. This process builds upon a generic cloud environment model (*CEM*). Besides the model elements for describing the *CECs*, the *CEM* contains further aspects relevant for providing migration support. Among those are elements for modeling the structure of specific cloud

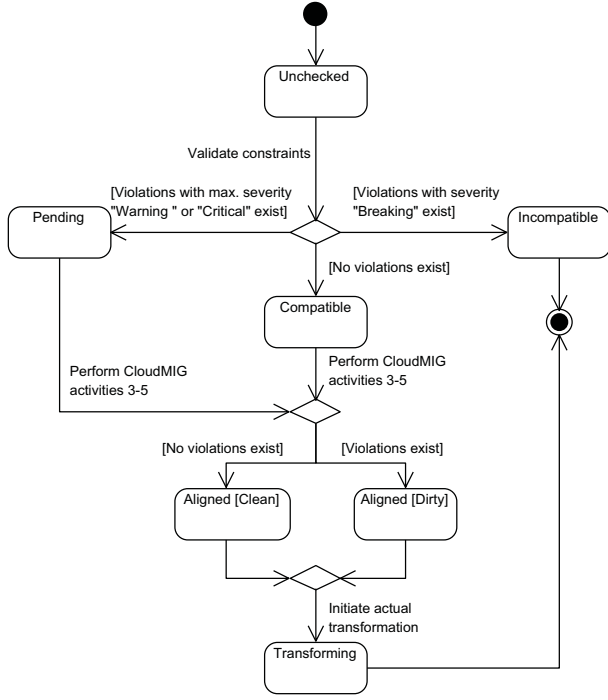


Figure 1. The possible states of a legacy software system for a specific target cloud environment in *CloudMIG* according to the validation of its *CECs*.

environments, legacy code data containers, or virtualized hardware resources, for instance. We further studied the approach's properties by modeling the PaaS environment Google App Engine for Java including the identified *CECs* and implemented the prototype tool *CloudMIG Xpress* that can execute the validation process. We report on a quantitative evaluation where we reconstructed the architectural models of five open source systems and investigated their intrinsic *CEC violations* by applying the validation process.

The remainder of the paper is structured as follows: Section II provides an overview of key aspects when migrating legacy software systems to the cloud with a focus on *CECs*. In Section III, we describe the extensible architecture for detecting *CEC violations*. Section IV presents the quantitative evaluation. The related work is described in Section V before Section VI draws the conclusions and outlines the future work.

II. MIGRATION TO THE CLOUD

A. Challenges

SaaS providers have to overcome numerous challenges and shortcomings of current approaches when migrating existing software systems to the cloud. Organizational implications might include a reshaping of internal divisions as responsibilities of IT or software maintenance departments shift, for instance. In addition, new liability or auditing issues may arise because sensible data is no longer stored exclusively on premise. Along with the question of which data assets can be moved to the cloud comes the increased need to encrypt this data. However, recent

advancements in the cryptography domain in achieving fully homomorphic encryption might eventually enable practicable arbitrary computation on encrypted data without the necessity to decrypt it beforehand and therefore mitigate data security concerns to a great extent [3].

Regarding technical challenges of a migration we identified the following major shortcomings of current approaches [2]: Solutions for migrating software systems to the cloud are limited to particular cloud providers (1). A poor level of automation concerning the creation of a target architecture, a mapping to this architecture, and the detection of *CEC violations* is prevalent (2). Further on, resource efficiency is not taken into account sufficiently. This is becoming particularly relevant regarding the prevailing pay-per-use billing models (3). Finally, automated support for evaluating a target architecture's scalability at design time is rare in the cloud computing context (4). These shortcomings constitute core drivers in the design of our overall approach *CloudMIG* being briefly described in the next Section II-B to explain the context. The definition of *CEM* and the quantitative evaluation regarding detected *CEC violations* presented in this paper address the shortcoming (1) and the last part of (2).

B. The Approach *CloudMIG*

CloudMIG is our approach for supporting reengineers to semi-automatically migrate existing software systems to cloud-based applications [2]. It incorporates usage patterns and varying resource demands in creating a target architecture candidate and concentrates on enterprise software systems offered by SaaS providers. The approach is composed of six major activities:

A1 Extraction: Includes the extraction of architectural and utilization models of the legacy system. They base upon OMG's Architecture-Driven Modernization¹ efforts and utilize the Knowledge Discovery Meta-Model² (KDM) and Software Metrics Meta-Model³ (SMM).

A2 Selection: Selection of an appropriate *CEM*-compatible cloud profile candidate.

A3 Generation: The generation activity produces the target architecture and a mapping model. Thereby, a model describing the target architecture's *CEC violations* is created that serves as an input artefact for the target architecture generation process. The violations are detected with the aid of the extensible architecture described in this paper.

A4 Adaptation: The adaptation activity enables a reengineer to manually adjust the target architecture.

A5 Evaluation: The evaluation activity involves static analyses and a runtime simulation of the target architecture.

A6 Transformation: The actual transformation of the existing system from the generated target architecture to the aimed cloud environment.

¹OMG ADM: <http://adm.omg.org/> (Accessed Dec. 21, 2010)

²KDM: <http://www.omg.org/spec/KDM/> (Accessed Dec. 21, 2010)

³SMM: <http://www.omg.org/spec/SMM/> (Accessed Dec. 21, 2010)




Figure 2. CloudMIG Xpress overview.

C. Cloud Environment Constraints

1) *Definition*: To clarify concepts, we explain the constraint-related key terms below that are used in this paper.

Cloud environment constraint (CEC): A constraint imposed by a cloud environment related to a specific technical action. When a (potential) guest application attempts to execute this specific action, the cloud environment prevents it from being executed.

Cloud environment constraint violation (CEC violation): Action of a (potential) guest application that would be prevented from being executed by a cloud environment due to a related CEC. The CEC violation is caused by an action that is manifested in a source code element of the guest application.

Violation Severity: The severity of a CEC violation indicates the likely effort to fix the CEC violation during a migration process. As this can vary widely depending on different legacy applications, we apply the *violation severity* rather pessimistic and propose three simple concrete severities: *Breaking*, *Critical*, and *Warning* associated with high, medium, and low effort, respectively.

It should be noted that the *violation severity* is biased according to the experience and subjective appraisal of a person modeling a cloud environment. Therefore, it should be rather seen as a hint for the reengineer. It is more important to detect a CEC violation at all and to make the reengineer beware of it.

2) Examples:

CEC: Using Google App Engine for Java, the total number of files is limited to 3,000 per default.

CEC violation: An application that exceeds this limit.

Violation severity: Warning (assuming that the creation of new container structures is a rather simple problem).

CEC: Using Amazon EC2, the local storage of VM instances is transient. For persistent storing one of Amazon's services like EBS, S3, or RDS has to be used.

CEC violation: An application that writes to the local file system in one of its methods.

Violation severity: Critical.

CEC: Using Google App Engine for Java, only JVM-compatible languages can be used for guest applications.

CEC violation: A C++ application.

Violation severity: Breaking.

3) *Management with CloudMIG*: CECs are made explicit in *CloudMIG*. Cloud profiles are modeled according to the CEM (see Section III-A) and describe a specific cloud environment in a reusable manner. The cloud profiles contain specifications of the CECs which the cloud environment establishes. Furthermore, *CloudMIG* provides means for automatically detecting CEC violations of a legacy application and points the reengineer to architectural elements that cause the violations. These detection mechanisms are designed to be extensible for cases where the standard constraint validators may not be sufficient.

On the one hand, the detected CEC violations are utilized by the reengineer to obtain a quick overview of problematic system parts which need special attention and as a basis for comparing competing cloud environment offers. On the other hand, the capability to detect CEC violations is used in the adaptation and evaluation activities (A4 and A5) to reason about the quality of different target architectures. As the reengineer has the possibility to manually adjust a target architecture, the detected CEC violations can vary over time and are therefore computed more than once.

Fig. 1 shows the integration of CECs in *CloudMIG*'s workflow. Before the initial validation of the CECs, a legacy software system is marked as "unchecked". Afterwards, further proceeding depends on the existence of CEC violations and the violation severities of detected CEC violations. If *Breaking* violations exist, the legacy system is considered as "incompatible". Referring to example three from Section II-C2 again, a migration would imply a transformation from one programming language into another. *CloudMIG* does not provide further support for such kinds of migrations and therefore the workflow ends in the case *Breaking* violations exist. Otherwise, the legacy system is either regarded as "compatible" (no CEC violations exist) or "pending" (*Critical* or *Warning* violations exist). The distinction is made as for "pending" legacy systems feedback for the reengineer and tracing

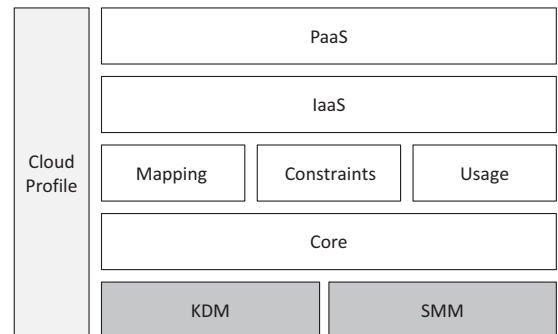


Figure 3. The packages of CEM.

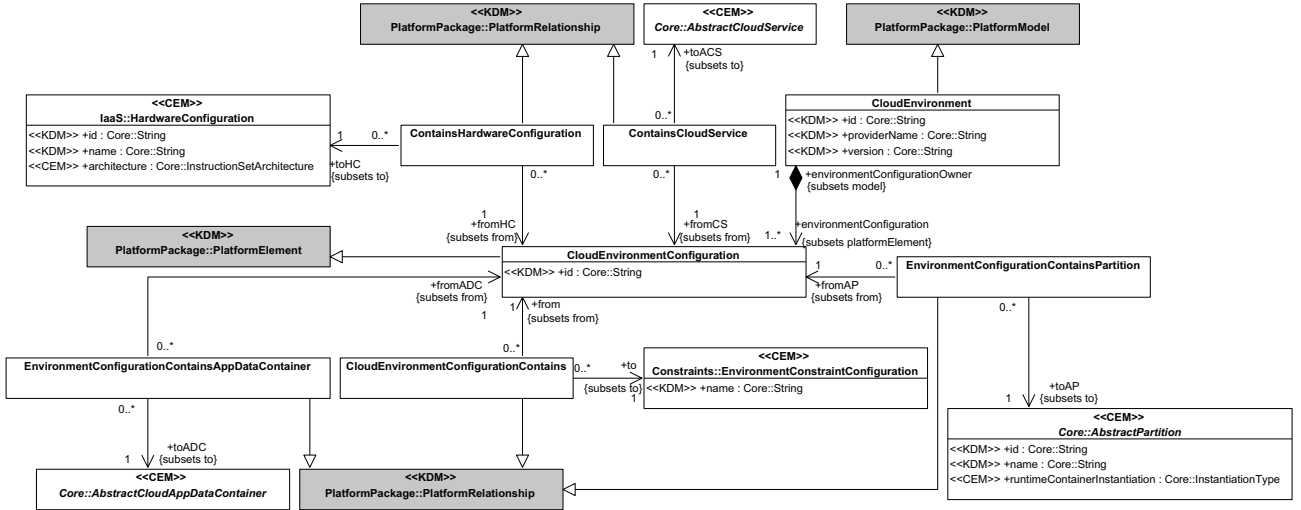


Figure 4. The “Cloud Profile” package of CEM (excerpt).

to the sources of the *CEC violations* has to be provided, for example. *CloudMIG*’s activities A3–A5 are executed subsequently in both cases. As mentioned above, the detected *CEC violations* can vary over time and therefore both states “Aligned (Clean)” (no *CEC violations* exist) and “Aligned (Dirty)” (*Critical* or *Warning* violations exist) can be reached along both paths after final constraint validation. Accomplishing the actual transformation ends the migration from *CloudMIG*’s perspective.

The *CECs* are analyzed in *CloudMIG* with so called constraint validators. Fig. 2 outlines our prototype tool *CloudMIG Xpress* and the integration of the constraint validators. As with components for most of the approach’s activities they can be plugged into the architecture, too. The components presented in Fig. 2 exhibit their affiliation to *CloudMIG*’s activities through related activity numbers.

III. THE ARCHITECTURE FOR DETECTING CEC VIOLATIONS

A. Cloud Environment Model

The Cloud Environment Model (*CEM*) constitutes the foundation for *CloudMIG*’s capabilities to detect *CEC violations*. It is aligned with OMG’s KDM and comprises a model for describing *CECs*. Instances of *CEM* represent specific cloud environments (cloud profiles) and have to be modeled only once and can then be reused by other reengineers. This applies to included *CECs* and already implemented constraint validators, too. It is intended to provide an appropriate public repository in the future.

1) *Overview*: The *CEM* is organized in layered packages as presented in Fig. 3. The *Core* package includes basic elements like abstract cloud services or partitions. The last allowing both Amazon EC2’s Availability Zones and Regions, for example. The further packages build upon the core package. The *Mapping* package comprises model elements that enable integration of legacy system parts

into a cloud environment. “Mapping” therefore means the assignment of legacy system parts to entities available in the cloud domain. Potential incompatibilities are handled by means of adapters that have to be created manually in the subsequent transformation step (*CloudMIG* activity A6), for instance. The *Usage Package* contributes model elements for describing and extracting the utilization model used by *CloudMIG*. In doing so, it incorporates measures and measurements modeled with OMG’s SMM. The *Constraints* package models the *CECs* and is covered in Section III-A3 in greater detail.

The *IaaS package* and *PaaS package* comprise elements for the corresponding cloud service models. The first follows the structural elements of the cross platform cloud API Deltacloud⁴ to some degree. The last is designed more generic as PaaS clouds exhibit an even broader bandwidth. The *Cloud Profile* package forms the entry point for modeling specific cloud environments. An excerpt is presented in Fig. 4. A *CloudEnvironment* can contain several *CloudEnvironmentConfigurations*. Taking Google App Engine as an example, there exists a *CloudEnvironmentConfiguration* for *Google App Engine for Java* and one for *Google App Engine for Python*. Furthermore, Fig. 4 shows elements for including an *IaaS HardwareConfiguration* (e.g. an Amazon EC2 “High-Memory Extra Large Instance”), an *EnvironmentConstraintConfiguration* for incorporating *CECs*, and the abstract classes *AbstractCloudService* and *AbstractCloudAppDataContainer* for providing convenient generic extension points for cases the concrete instances in other layers may not be sufficient. The alignment with KDM is described in the following Section.

2) *Alignment with KDM*: Besides domain-specific elements of the *Cloud Profile* package, one can recognize

⁴<http://www.deltacloud.org/> (Accessed Dec. 21, 2010)

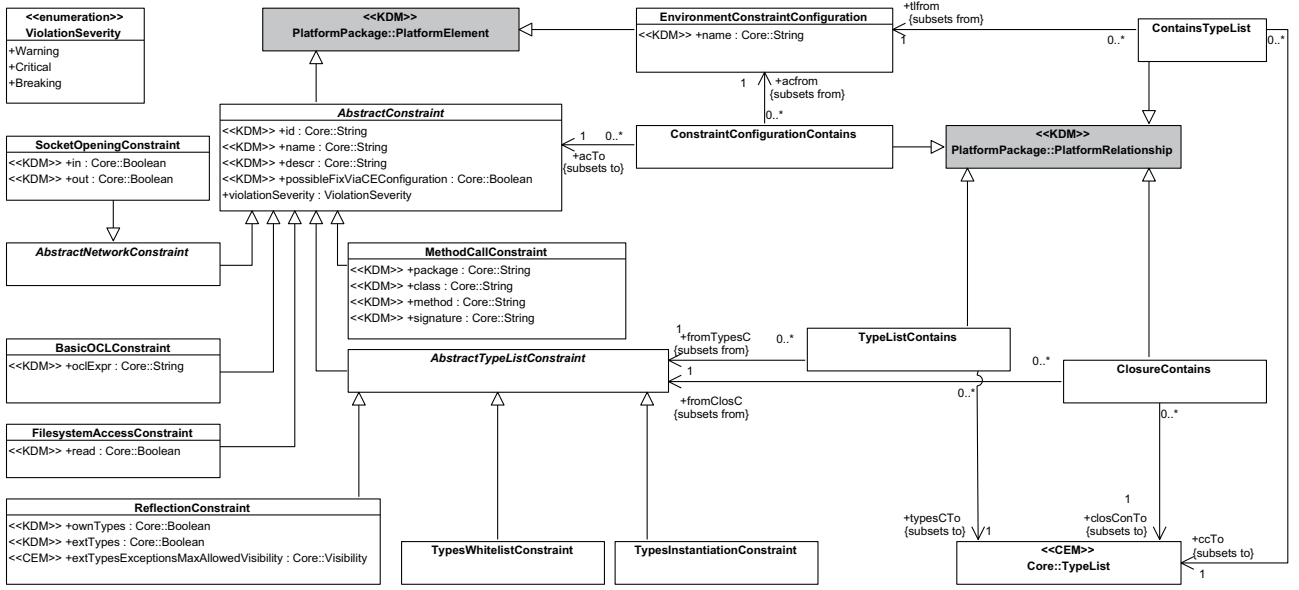


Figure 5. The “Constraints” package of CEM (excerpt).

```

1 <extension name="cloudprofile">
2   <stereotype name="CloudEnvironment" type="
     PlatformModel">
3     <tag tag="id" type="String"/>
4     <tag tag="providerName" type="String"/>
5     <tag tag="version" type="String"/>
6   </stereotype>
7   <stereotype name="CloudEnvironmentConfiguration"
     type="PlatformElement">
8     <tag tag="id" type="String"/>
9   </stereotype>
10 </extension>
11
12 <model xsi:type="platform:PlatformModel"
     stereotype="// @extension.0/ @stereotype.0">
13   <taggedValue xsi:type="kdm:TaggedValue"
     tag="// @extension.0/ @stereotype.0/ @tag.0"
14     value="cloudmig.cloudprofiles.gae"/>
15   <platformElement xsi:type="
     platform:PlatformElement" stereotype="//
     @extension.0/ @stereotype.1">
16     <taggedValue xsi:type="kdm:TaggedValue" tag="
     // @extension.0/ @stereotype.1/ @tag.0"
     value="cloudmig.cloudprofiles.gae.java"/>
17   </platformElement>
18   <ownedRelation xsi:type="
     platform:PlatformRelationship" to="//
     @model.0/ @platformElement.1" from="//
     @model.0/ @platformElement.0"/>
19 </model>
20 </model>

```

Listing 1. Google App Engine for Java Cloud Profile (KDM excerpt).

the alignment of *CEM* with KDM in Fig. 4. Generally, *CEM* builds upon KDM’s platform and structure packages following the piggyback pattern [4] for the realization of DSLs. We provide a transformation from the domain model implemented as an Ecore model to a KDM-compatible version. Future KDM-conform modernization tools may therefore be able to process our model. The compatibility is achieved by avoiding to introduce new meta model elements. The *CEM* classes inheriting from KDM classes (gray) in Fig. 4 are rather transformed

to KDM Stereotypes, their attributes become KDM TagDefinitions, and the *CEM* packages are realized as KDM ExtensionFamilies, to cover just the major modeling elements. This approach uses the light-weight extension mechanism of KDM. The mentioned elements constitute new so called virtual meta-model elements. They have to consider given restrictions. For example, the *CloudEnvironmentConfiguration* inherits from a KDM *PlatformElement* and therefore can not incorporate the *EnvironmentConstraintConfiguration* class (a *PlatformElement* as well, see Fig. 5) via a composition, as this association does not exist for two *PlatformElements*. It rather has to utilize an instance of KDM’s *PlatformRelationship* (the class *CloudEnvironmentConfigurationContains*). Listing 1 illustrates this concept by means of a KDM Google App Engine for Java cloud profile extract in XMI notation. For example, the *PlatformModel* in line 12 represents CEM’s *CloudEnvironment*, as its stereotype attribute refers to the according KDM Stereotype in line 2. The virtual meta-model element *CloudEnvironment* is in turn contained in a KDM *ExtensionFamily*⁵ describing *CEM*’s *Cloud Profile package* (line 1).

3) *Modeling of Constraints*: *CEM*’s *EnvironmentConstraintConfiguration* class was already mentioned in the former Section III-A2. It is located in the *CEM Constraints package*. An excerpt of this package is presented in Fig. 5. An *EnvironmentConstraintConfiguration* constitutes a container for a set of *CECs*.

⁵The KDM Ecore model used from the tool MoDisco names the “extensionFamily” role of the according containment relationship in the “KDMFramework” super class element merely “extension”

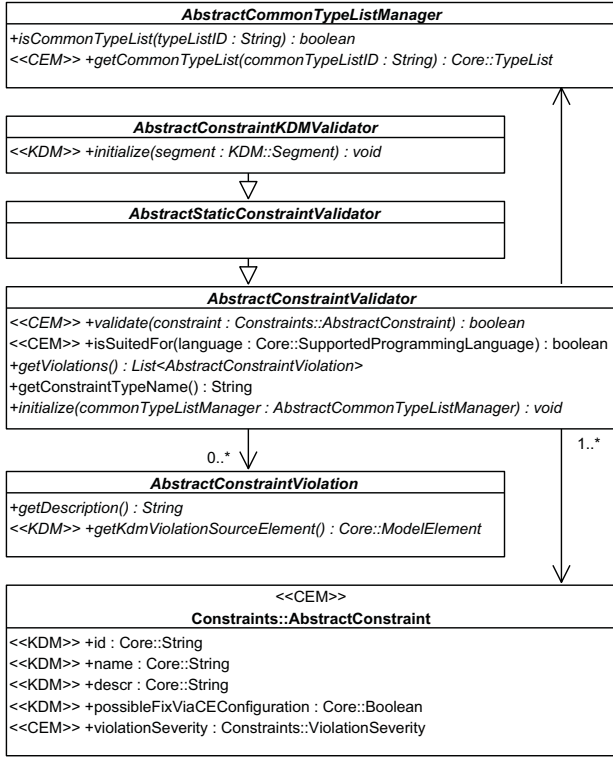


Figure 6. Classes incorporated in the CEC validation process.

Concrete CECs are modeled by inheriting from `AbstractConstraint` which provides an attribute for a `ViolationSeverity`, for instance. In Fig. 5, some concrete CEC examples are listed for illustrating purposes, the `FileSystemAccessConstraint` for detecting a forbidden file system access, or the `ReflectionConstraint` for discovering a prohibited call to the reflection API, for example. The `BasicOCLConstraint` allows a generic detection on MOF-compliant models like KDM. Instances of `TypeList` provide a convenient way to refer to a predefined set of types, or to create an own set that can be used for modeling CECs. For example, *CloudMIG Xpress* contains a predefined set of qualified JRE 6 types that can be applied in Google App Engine for Java's case in a `TypesWhitelistConstraint` as a superset (`ClosureContains` relationship), as Google App Engine for Java restricts access to only a subset of these types (`TypeListContains` relationship).

B. Constraint Validation

Classes playing a role in the CEC validation process are shown in Fig. 6. Constraint validation plugins for *CloudMIG Xpress* (see Fig. 2) have to provide a subclass of `AbstractConstraintValidator`. Currently, there exist two further abstract subclasses one can inherit from for convenience reasons, namely variants for static validation (`AbstractStaticConstraintValidator`) and

for static validation with respect to KDM source models (`AbstractConstraintKDMValidator`). Constraint validators refer to a specific type of CEC, for a given CEC there may exist various constraint validators. An `AbstractCommonTypeListManager` handles the predefined lists of types mentioned in Section III-A3.

The CEC validation process basically functions as follows. *CloudMIG Xpress* asks every present constraint validator whether the programming language of the source model is supported (`isSuitedFor`-method) and if the offered validation capability matches a constraint type present in a cloud profile (`getConstraintTypeName`-method). It should be noted that in the case for language independent validation plugins the first method always returns `true`. If an appropriate constraint validator is found, *CloudMIG Xpress* initializes it and launches the specific validation process with a call to the method `initialize` and `validate`, respectively. In the case CEC violations are detected, the constraint validator returns a list of `AbstractConstraintViolations` due to a call to the `getViolations` method after the plugin finished validation.

IV. EVALUATION

To study characteristics of CEC violation detection and to achieve insights in quantity, types, and properties of CEC violations a reengineer might face when considering a migration of an existing system to the cloud, we evaluated our approach by means of five open source systems and Google App Engine for Java.⁶ Our prototype tool *CloudMIG Xpress* was utilized to process the reconstructed models of these systems and to perform the validation process to detect CEC violations. An exemplary summary of a single validation run can be seen in Fig. 7. In the context of our evaluation, for each CEC violation detected a data record was reported comprising the IDs of the processed application, the particular KDM model, and the class referring to the source of the CEC violation, as well as the corresponding violation severity and the type of the constraint.

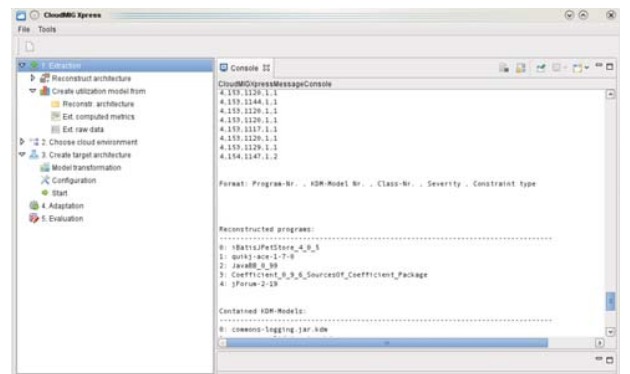


Figure 7. CloudMIG Xpress displaying a constraint validation summary.

⁶<https://code.google.com/intl/en-EN/appengine/docs/java/overview.html> (Accessed Dec. 21, 2010)




Figure 8. Amount of classes causing constraint violations (w/o third-party libs).

A. Background

We modeled the PaaS Cloud Environment Google App Engine for Java to evaluate our approach. Accordingly, we created a *CEM* instance comprising structural elements and the corresponding *CECs*. In *CEM*'s terms, Google App Engine for Java constitutes a *CloudEnvironmentConfiguration* as part of the *CloudEnvironment* Google App Engine, as mentioned in Section III-A1 before. It allows JVM-compatible languages to be used by guest applications and provides them an isolated sandbox environment with a considerable number of restrictions. Despite not all restrictions are documented extensively, they form a sufficiently well-suited basis for modeling the related *CECs*. The information concerning functionality, structure, and sandbox restrictions were distilled from Google's webpage. Moreover, several further web logs provided helpful information for understanding more facets of the specific *CECs*.

In the context of this evaluation, an important aspect was judging about the feasibility of our approach and less exploring its applicability on a wide range of diverging system types. Therefore, we narrowed down the type of potential applications to Java and web-based systems. The following open source applications were selected for evaluation purposes:

- App1: Coefficient Core⁷ V. 0.9.6
- App2: iBATIS JPetStore⁸ V. 4.0.5
- App3: JavaBB⁹ V. 0.99
- App4: jForum¹⁰ V. 2.1.9
- App5: Ace Operator¹¹ V. 1.7.0

⁷<http://coefficient.sourceforge.net/> (Accessed Dec. 21, 2010)

⁸<http://ibatisjpetstore.sourceforge.net/> (Accessed Dec. 21, 2010)

⁹<http://www.javabb.org/> (Accessed Dec. 21, 2010)

¹⁰<http://jforum.net/> (Accessed Dec. 21, 2010)

¹¹http://www.quik-j.com/ace_operator.htm (Accessed Dec. 21, 2010)

Table I
THE OPEN SOURCE SYSTEMS UTILIZED IN THE EVALUATION.

App	Name	Domain	#Classes (w/o libs)	#Libraries	LOC (w/o libs)
1	Coefficient Core 0.9.6	Collaboration platform	131	41	11,862
2	iBATIS JPetStore 4.0.5	Pet store	43	12	2,132
3	JavaBB 0.99	Forum software	256	43	12,239
4	jForum 2.1.9	Forum software	316	30	29,563
5	Ace Operator 1.7.0	Live support	556	26	69,516

Some basic characteristics of each system can be found in table I. Considering LOC of application's sources while leaving out used third-party libraries, the sizes of the largest and the smallest application differ by a factor of approximately 33 and therefore allow analyses for varying scales.

B. Methodology

For extracting the related KDM models of the applications, we applied the tool MoDisco V. 0.8.¹² In our analysis, we distinguished cases where we solely considered the application's own sources and those taking into account third-party libraries. This is due to the fact that some needed elements necessary for certain constraint validators were not present in the KDM models for the libraries discovered with MoDisco.

For example, the models lacked instances of the class "Calls" representing method calls among others. 29 *CECs* were modeled, it was possible to detect 20 *CECs* that could be covered by our provided validators corresponding to 8 covered types of *CECs*.




Figure 9. The detected constraint violation densities per application (w/o third-party libs).

¹²<http://www.eclipse.org/MoDisco/> (Accessed Dec. 21, 2010)




Figure 10. Distribution of detected violation types (w/o third-party libs).

The lack of the residual *CECs* is caused by the fact that *CloudMIG Xpress* does, at the moment, not provide support for detecting violations that are only identifiable at runtime and just static analyses are implemented.

Table II lists the types of *CECs* that could be detected during the evaluation and the nr. of present variants. For example, several variants of a *MethodCallConstraint* exist due to restrictions on calls to different methods. It should be noted that for different variants of a *CEC* type there can be dissimilar *violation severities* assigned. Furthermore, the *CECs* inheriting from *AbstractTypeListConstraint* show only a single variant in the table. But looking at a *TypesWhitelistConstraint*, it translates to a prohibition to access 2,388 of the JRE types, for example.

We examined the following questions of special interest:

- Q1** How many classes of an application raise *CEC violations*?
- Q2** How does the density of *CEC violations* vary among applications?
- Q3** What types of *CEC violations* are prevalent?
- Q4** Regarding classes that raise *CEC violations*, how does the size of such classes relate to the number of *CEC violations* they raise?
- Q5** How does the density of *CEC violations* vary among classes?
- Q6** How does the number of *CEC violations* raised by an application's own sources relate to the number of *CEC violations* raised by utilized third-party libraries?

Additionally, for every question stated above we were interested whether we could identify indicators for patterns regarding a correlation with the systems' sizes.

The number of classes and libraries for each program could be extracted from the KDM models. As a further sufficiently well-suited size measure we employed LOC. It was measured with CLOC¹³ V. 1.52 omitting comment and blank lines. However, it was only possible to measure LOC for the applications' own sources, as the sources for the third-party libraries were missing. However, third-party libraries were only incorporated for analyzing *Q6* and for that question LOC is not relevant.

C. Results

The absolute numbers of detected *CEC violations* differ widely. For App1-App5, there were (87/4,386), (3/8), (98/932), (273/1,428), (7,795/612) *CEC violations* detected (own sources/ third-party libraries).

Q1: Fig. 8 shows the amount of classes causing *CEC violations*. The range spans from approx. 2.5% to 33% for the smallest (App2) and the largest application (App5), respectively. Generally, with size comes a steady growth in the nr. of classes raising *CEC violations*, but without a uniform growth rate or pattern.

Q2: Fig. 9 shows the detected constraint violation densities per application. They are similar for App1, App3, and App4 considering *Critical violation severities*, despite App4 being roughly 3 times bigger than the others. The smallest application App2's density is slightly lower, but the biggest application App5's density exceeds the others in orders of magnitude. Generally, we conclude that the size can also be an indicator for higher densities of *CEC*

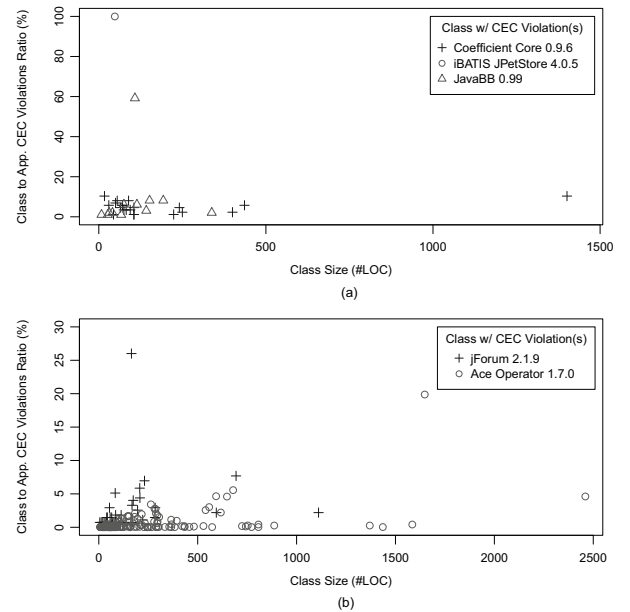


Figure 11. Classes being responsible for an amount of constraint violations and the relation to their size (w/o third-party libs). Split into two diagrams (a) and (b) due to legibility.

¹³<http://cloc.sourceforge.net/> (Accessed Dec. 21, 2010)




Figure 12. The detected constraint violation densities per class (w/o third-party libs).

violations, too. But for *CEC violations* with *Warning* as *violation severity* this statement is weaker.

Q3: In Fig. 10 the distribution of detected violation types is presented. It is no surprise that the *TypesWhitelistConstraint* is dominant, as this *CEC* can match for a plethora of cases as stated in Section IV-B. APP2 does not coincide, but this is likely due to the low overall number of *CEC violations* found for this application. *CEC violations* related to restricted file system access are in addition to it quite frequently observed.

Q4: Fig. 11 shows a scatterplot for classes raising *CEC violations* and their relation to the relative number of *CEC violations* they raise and their size. The predominant number of those classes are responsible for 5% or less of the overall number of *CEC violations* raised. The root causes are wide spread over the systems. But there exist some outliers, the likely most spectacular one referring to App3 and being responsible for approx. 60% of *CEC violations*. All *CEC violations* from the small App2 are located in one class.

Q5: Fig. 12 shows the detected constraint violation densities per class. The largest application App5 exhibits the largest jitter. It is remarkable that the medians of all classes in all applications lie in a rather narrow band of approx. 3-5 (normalized to 100 LOC for each class).

Q6: In Fig. 13 the third-party libraries are incorporated. It presents the origin of *CEC violations*. Already the total numbers stated above showed vast differences. Once more, App5 is a special case, as it is the only system that

Table II
THE INCORPORATED *CECs*.

Type	#Variants	Violation severity
MaxTotalNrOffFilesConstraint	1	1 Warning
MethodCallConstraint	12	3 Critical, 9 Warning
SocketOpeningConstraint	1	1 Critical
FilesystemAccessConstraint	2	1 Critical, 1 Warning
ReflectionConstraint	1	1 Critical
TypesInstantiationConstraint	1	1 Critical
TypesWhitelistConstraint	1	1 Critical
LanguageConstraint	1	1 Breaking

produces substantially more *CEC violations* in its sources than its third-party libraries do. Furthermore, it can be seen in Fig. 13 a)-c) and partially in d) that few libraries are responsible for most of the *CEC violations*. Identifying and handling those primarily seems to be a worthwhile approach.

D. Discussion

The evaluation investigates *CEC violations* in a PaaS context. Compared to current IaaS environments the individual PaaS offerings are more heterogeneous as they add diverse abstraction layers or specific preconfigured software stacks to the basic infrastructure building blocks. From a cloud user perspective these additional boundaries constitute further *CECs* that need to be taken into account. Considering other PaaS environments, the results from the evaluation will primarily shift because of two factors. First, PaaS environments that do not support the Java runtime environment will yield incomparable results. As stated in Section II-C3 the approach *CloudMIG* does not aim to migrate software systems through applying a programming language transformation. It terminates its workflow if *CEC violations* with the assigned *violation severity Breaking* are detected. The *LanguageConstraint* listed in table II is a *Breaking* constraint. Second, the majority of detected *CEC violations* are *TypesWhitelistConstraints*.




Figure 13. Origin of constraint violations, light gray slices represent third-party libraries.

Hence, concordance of results for other PaaS environments offering a JRE compatibility will particularly be sensitive to type restrictions diverging from those defined in Google App Engine for Java's sandbox environment.

Regarding IaaS environments, it can be expected that the number of detected violations would be considerably lower. This is due to the fact that IaaS environments by definition lack narrow restrictions concerning the underlying software stack that in turn could translate into additional modeled *CECs* and violations of those.

Furthermore, identifying orphaned classes, components, or even subsystems is a typical task in reengineering projects. Here, static and dynamic analyses can be used. Considering only system parts that are actually being used might also lead to a significant reduction of *CEC violations*. This is planned for future analyses.

Moreover, only system types were analyzed that were relatively well-suited. First, our approach assumes that KDM models can be extracted. This might not be the case for many existing programming languages as they currently lack appropriate tool support. As mentioned before, the extracted models from third-party libraries were also incomplete and therefore results for *Q6* can rather show a tendency. However, the used tool MoDisco provides a suitable framework for developing and enhancing discoverers. Second, Google App Engine for Java supports web-based Java software systems. Only representatives of this application type were studied. The number of probands was also rather small.

Overall, generalizability is limited but the experiments provided us valuable insights and showed that our approach can be successfully utilized.

V. RELATED WORK

The reengineering of software systems builds a primary domain our approach is located in. An overview of legacy software system migration was contributed in [5]. The migration of existing systems to new platforms shows inherent complexity and numerous difficulties a reengineer has to overcome and migrating legacy systems to cloud-based environments makes no difference at this point. The authors in [6] propose several techniques to reduce migration complexity, for example dynamic program analysis, software visualization, and knowledge discovery.

The novel field of cloud computing constitutes a major foundation of our work as well. [7] provides an overview of 20 cloud definitions and extracts a consensus definition. The authors in [8] survey the research published in this area and discuss lessons learned from related technologies. A related case study was conducted in [9]. The authors report on a migration that transferred an enterprise software system to an IaaS cloud environment. However, the case study concentrates on financial and socio-technical enterprise issues. [10] reports on a migration of a large scientific database to the cloud.

VI. CONCLUSION AND FUTURE WORK

We presented our extensible architecture for detecting a legacy software system's violations of constraints imposed by cloud environments when considering a migration. It is an important constituent of our approach *CloudMIG* which supports reengineers in migrating existing software systems to the cloud. Here, the detection and highlighting of crucial system parts is an essential early phase activity. Violations may be easy to fix, but a reengineer has to be aware of them. Furthermore, our approach provides support for assessing the severity of a detected violation.

The future work includes further analyses of constraints impeding a migration to the cloud. Moreover, it is planned to model additional cloud profiles, provide a public repository, and to improve detection capabilities. For example, incorporating runtime information constitutes a promising approach. Furthermore, it will be interesting to examine application types that differ to a greater extent.

REFERENCES

- [1] M. Armbrust, A. Fox, R. Griffith, A. D. Joseph, R. H. Katz, A. Konwinski, G. Lee, D. A. Patterson, A. Rabkin, I. Stoica, and M. Zaharia, "Above the Clouds: A Berkeley View of Cloud Computing," EECS Department, University of California, Berkeley, Tech. Rep. UCB/EECS-2009-28, Feb 2009.
- [2] S. Frey and W. Hasselbring, "Model-Based Migration of Legacy Software Systems to Scalable and Resource-Efficient Cloud-Based Applications: The CloudMIG Approach," in *Cloud Computing 2010: Proceedings of the 1st International Conference on Cloud Computing, GRIDs, and Virtualization*, 2010, (to appear).
- [3] C. Gentry, "Computing Arbitrary Functions of Encrypted Data," *Commun. ACM*, vol. 53, no. 3, pp. 97–105, 2010.
- [4] D. Spinellis, "Notable design patterns for domain-specific languages," *Journal of Systems and Software*, vol. 56, no. 1, pp. 91–99, 2001.
- [5] J. Bisbal, D. Lawless, B. Wu, J. Grimson, V. Wade, R. Richardson, and D. O'Sullivan, "An overview of legacy information system migration," dec. 1997, pp. 529–530.
- [6] L. Wu, H. Sahraoui, and P. Valtchev, "Coping with legacy system migration complexity," in *Proceedings. 10th IEEE International Conference on Engineering of Complex Computer Systems, 2005. ICECCS 2005.*, jun. 2005, pp. 600–609.
- [7] L. M. Vaquero, L. Roderio-Merino, J. Caceres, and M. Lindner, "A break in the clouds: towards a cloud definition," *SIGCOMM Comput. Commun. Rev.*, vol. 39, no. 1, pp. 50–55, 2009.
- [8] I. Sriram and A. Khajeh-Hosseini, "Research Agenda in Cloud Technologies," *CoRR*, vol. abs/1001.3259, 2010.
- [9] A. Khajeh-Hosseini, D. Greenwood, and I. Sommerville, "Cloud Migration: A Case Study of Migrating an Enterprise IT System to IaaS," *CoRR*, vol. abs/1002.3492, 2010.
- [10] A. Thakar and A. Szalay, "Migrating a (Large) Science Database to the Cloud," in *HPDC '10: Proceedings of the 19th ACM International Symposium on High Performance Distributed Computing*. New York, NY, USA: ACM, 2010, pp. 430–434.